

Timestep sensitivity analysis: Exponential vs Euler discretisation

Sandra Montes (@slmontes)

2026-01-20

Introduction

This script performs a timestep sensitivity analysis for Case Scenario 1 (lockdown) comparing two discretisation methods:

Method 1: Exponential discretisation

- $\text{infection}[t] == (1 - \exp(-(1 - u[t]) * \beta * I[t] * dt)) * S[t]$
- $\text{recovery}[t] == (1 - \exp(-\gamma * dt)) * I[t]$

Method 2: Simple Euler discretisation

- $\text{infection}[t] == (1 - u[t]) * \beta * I[t] * dt * S[t]$
- $\text{recovery}[t] == \gamma * dt * I[t]$

```
using JuMP
using Ipopt
using Plots
Plots.default(fmt = :png)    # or :svg
using DataFrames
using Statistics
```

We add a warmup step to trigger JIT compilation and eliminate compilation overhead time

```
function warmup_jit()
    warmup_model = Model(Ipopt.Optimizer)
    set_optimizer_attribute(warmup_model, "print_level", 0) # Silent

    # Small test problem
```

```

@variable(warmup_model, 0 <= x[1:5] <= 1)
@variable(warmup_model, y[1:4] >= 0)
@constraint(warmup_model, [t=1:4], y[t] == (1 - exp(-0.5 * x[t])) * x[t])
@constraint(warmup_model, [t=1:4], x[t+1] == x[t] - y[t])
@objective(warmup_model, Min, x[5])

optimize!(warmup_model)
end
warmup_jit()

```

Model parameters

```

β = 0.5          # infection rate
γ = 0.25         # recovery rate
u_max = 0.5      # maximum intervention
u_total = 10.0   # maximum cost

t0 = 0.0
tf = 100.0

S0 = 0.99
I0 = 0.01
C0 = 0.00;

```

```

# Timesteps to test
dt_values = [2.0, 1.0, 0.5, 0.25, 0.1, 0.05]

# Number of repetitions per dt cases (to calculate mean and std)
n_repetitions = 50

# Initialise results storage
results_exponential = []
results_euler = [];

```

Running optimisation for each timestep on both methods:

Method 1: Exponential discretisation

```

for dt in dt_values

    # Number of time points
    T = Int(tf/dt)

    # Storage results
    solve_times_exp = []
    objectives_exp = []
    statuses_exp = []

    # Trajectories
    local v_opt_exp, C_opt_exp, S_opt_exp, I_opt_exp

    # Run n_repetitions times
    for rep in 1:n_repetitions
        model_exp = Model(Ipopt.Optimizer)
        set_optimizer_attribute(model_exp, "print_level", 0)

        # Variables
        @variable(model_exp, 0 <= S[1:(T+1)] <= 1)
        @variable(model_exp, 0 <= I[1:(T+1)] <= 1)
        @variable(model_exp, 0 <= C[1:(T+1)] <= 1)
        @variable(model_exp, 0 <= u[1:(T+1)] <= u_max)

        # Exponential discretisation
        @expressions(model_exp, begin
            infection[t in 1:T], (1 - exp(-(1 - u[t]) * β * I[t] * dt)) * S[t]
            recovery[t in 1:T], (1 - exp(-γ * dt)) * I[t]
        end)

        @constraints(model_exp, begin
            S[1]==S0
            I[1]==I0
            C[1]==C0
            [t=1:T], S[t+1] == S[t] - infection[t]
            [t=1:T], I[t+1] == I[t] + infection[t] - recovery[t]
            [t=1:T], C[t+1] == C[t] + infection[t]
            dt * sum(u[t] for t in 1:T+1) <= u_total
        end)

        # Objective

```

```

@objective(model_exp, Min, C[T+1])

# Solve and time
solve_start = time()
optimize!(model_exp)
solve_time = time() - solve_start

# Extract results
status = termination_status(model_exp)
obj_value = objective_value(model_exp)

# Store results
push!(solve_times_exp, solve_time)
push!(objectives_exp, obj_value)
push!(statuses_exp, status)

# Save trajectories from first run
if rep == 1
    u_opt_exp = value.(u)
    C_opt_exp = value.(C)
    S_opt_exp = value.(S)
    I_opt_exp = value.(I)
end
end

# Results
mean_time_exp = mean(solve_times_exp)
std_time_exp = std(solve_times_exp)

# Verify ALL runs converged successfully
if !all(s == MOI.LOCALLY_SOLVED for s in statuses_exp)
    error("Not all Exponential runs converged! Statuses: $statuses_exp")
end
status_exp = statuses_exp[1] # All LOCALLY_SOLVED

# Verify that all objectives are identical
if !all(abs(obj - objectives_exp[1]) < 1e-10 for obj in objectives_exp)
    error("Objectives varied across repetitions:\n" *
        " Values: $objectives_exp\n" *
        " Max difference: $(maximum(objectives_exp) - minimum(objectives_exp)))
end
obj_value_exp = objectives_exp[1]

```

```

# Control
threshold = 0.01
control_start_idx = findfirst(x -> x > threshold, u_opt_exp)
control_start_exp = control_start_idx !== nothing ? (control_start_idx - 1)
control_active = [u_opt_exp[i] > threshold for i in 1:length(u_opt_exp)]
control_duration_exp = sum(control_active) * dt

# Exponential results
result_exp = Dict(
    :dt => dt,
    :method => "Exponential",
    :time_points => T+1,
    :objective => obj_value_exp,
    :solve_time => mean_time_exp,
    :solve_time_std => std_time_exp,
    :solve_times_all => solve_times_exp,
    :convergence => string(status_exp),
    :control_start => control_start_exp,
    :control_duration => control_duration_exp,
    :final_infections => C_opt_exp[end],
    :u_trajectory => u_opt_exp,
    :C_trajectory => C_opt_exp,
    :S_trajectory => S_opt_exp,
    :I_trajectory => I_opt_exp
)
push!(results_exponential, result_exp)
end

```

Method 2: Simple Euler discretisation (Linear approximation)

```

for dt in dt_values

    # Number of time points
    T = Int(tf/dt)

    # Storage results
    solve_times_euler = []
    objectives_euler = []
    statuses_euler = []

```

```

# Trajectories
local v_opt_euler, C_opt_euler, S_opt_euler, I_opt_euler

# Run n_repetitions times
for rep in 1:n_repetitions
    model_euler = Model(Ipopt.Optimizer)
    set_optimizer_attribute(model_euler, "print_level", 0)

    # Variables
    @variable(model_euler, 0 <= S[1:(T+1)] <= 1)
    @variable(model_euler, 0 <= I[1:(T+1)] <= 1)
    @variable(model_euler, 0 <= C[1:(T+1)] <= 1)
    @variable(model_euler, 0 <= u[1:(T+1)] <= u_max)

    # Simple Euler discretisation
    @expressions(model_euler, begin
        infection[t in 1:T], (1 - u[t]) * β * I[t] * dt * S[t]
        recovery[t in 1:T], γ * dt * I[t]
    end)

    @constraints(model_euler, begin
        S[1]==S0
        I[1]==I0
        C[1]==C0
        [t=1:T], S[t+1] == S[t] - infection[t]
        [t=1:T], I[t+1] == I[t] + infection[t] - recovery[t]
        [t=1:T], C[t+1] == C[t] + infection[t]
        dt * sum(u[t] for t in 1:T+1) <= u_total
    end)

    # Objective
    @objective(model_euler, Min, C[T+1])

    # Solve and time
    solve_start = time()
    optimize!(model_euler)
    solve_time = time() - solve_start

    # Extract results
    status = termination_status(model_euler)
    obj_value = objective_value(model_euler)

```

```

# Store results
push!(solve_times_euler, solve_time)
push!(objectives_euler, obj_value)
push!(statuses_euler, status)

# Save trajectories from first run
if rep == 1
    u_opt_euler = value.(u)
    C_opt_euler = value.(C)
    S_opt_euler = value.(S)
    I_opt_euler = value.(I)
end
end

# Calculate statistics
mean_time_euler = mean(solve_times_euler)
std_time_euler = std(solve_times_euler)

# Verify ALL runs converged successfully
if !all(s == MOI.LOCALLY_SOLVED for s in statuses_euler)
    error("Not all Euler runs converged! Statuses: $statuses_euler")
end
status_euler = statuses_euler[1]

# Verify all objectives are identical
if !all(abs(obj - objectives_euler[1]) < 1e-10 for obj in objectives_euler)
    error("Objectives varied across repetitions:\n" *
          " Values: $objectives_euler\n" *
          " Max difference: $(maximum(objectives_euler) - minimum(objectives_euler))")
end
obj_value_euler = objectives_euler[1]

# Control
threshold = 0.01
control_start_idx = findfirst(x -> x > threshold, u_opt_euler)
control_start_euler = control_start_idx !== nothing ? (control_start_idx - 1):length(u_opt_euler)
control_active = [u_opt_euler[i] > threshold for i in 1:length(u_opt_euler)]
control_duration_euler = sum(control_active) * dt

# Verify all objectives are identical
if !all(abs(obj - objectives_euler[1]) < 1e-10 for obj in objectives_euler)
    error("Euler: Not deterministic! Objectives varied across repetitions:\n$objectives_euler")
end

```

```

        "  Values: $objectives_euler\n" *
        "  Max difference: $(maximum(objectives_euler) - minimum(objectives_euler))
    end

    # Euler results
    result_euler = Dict(
        :dt => dt,
        :method => "Euler",
        :time_points => T+1,
        :objective => obj_value_euler,
        :solve_time => mean_time_euler,
        :solve_time_std => std_time_euler,
        :solve_times_all => solve_times_euler,
        :convergence => string(status_euler),
        :control_start => control_start_euler,
        :control_duration => control_duration_euler,
        :final_infections => C_opt_euler[end],
        :u_trajectory => u_opt_euler,
        :C_trajectory => C_opt_euler,
        :S_trajectory => S_opt_euler,
        :I_trajectory => I_opt_euler
    )
    push!(results_euler, result_euler)
end

```

Exponential discretisation results:

6×9 DataFrame							
Row	dt	time_points	objective	solve_time_mean	solve_time_std	converge	
	Float64	Float64	Float64	Float64	Float64	String	
1	2.0	51.0	0.731	0.0228	0.0053	LOCALLY	
2	1.0	101.0	0.6686	0.05	0.0033	LOCALLY	
3	0.5	201.0	0.6332	0.2312	0.0053	LOCALLY	
4	0.25	401.0	0.6145	0.2603	0.0123	LOCALLY	
5	0.1	1001.0	0.6028	1.273	0.1185	LOCALLY	
6	0.05	2001.0	0.5989	1.8665	0.1421	LOCALLY	

Simple Euler discretisation results:

6×9 DataFrame

Row	dt	time_points	objective	solve_time_mean	solve_time_std	converg
	Float64	Float64	Float64	Float64	Float64	String
1	2.0	51.0	0.5853	0.0312	0.0078	LOCALLY
2	1.0	101.0	0.5905	0.0466	0.0244	LOCALLY
3	0.5	201.0	0.5928	0.2013	0.008	LOCALLY
4	0.25	401.0	0.5939	0.3019	0.0319	LOCALLY
5	0.1	1001.0	0.5945	1.4404	0.0862	LOCALLY
6	0.05	2001.0	0.5947	1.3914	0.0918	LOCALLY

Relative errors of each method using 'dt=0.05' as reference

```

ref_idx_exp = findfirst(r -> r[:dt] == 0.05, results_exponential)
ref_idx_euler = findfirst(r -> r[:dt] == 0.05, results_euler)

if ref_idx_exp !== nothing && ref_idx_euler !== nothing
    ref_obj_exp = results_exponential[ref_idx_exp][:objective]
    ref_obj_euler = results_euler[ref_idx_euler][:objective]

    # Exponential method
    rel_errors_exp = []
    for r in results_exponential
        obj_error = abs(r[:objective] - ref_obj_exp) / ref_obj_exp * 100
        push!(rel_errors_exp, Dict(
            :dt => r[:dt],
            :obj_rel_error_pct => obj_error,
            :method => "Exponential"
        ))
    end

    # Euler method
    rel_errors_euler = []
    for r in results_euler
        obj_error = abs(r[:objective] - ref_obj_euler) / ref_obj_euler * 100
        push!(rel_errors_euler, Dict(
            :dt => r[:dt],
            :obj_rel_error_pct => obj_error,
            :method => "Euler"
        ))
    end

    errors_exp_df = DataFrame(rel_errors_exp)

```

```

errors_euler_df = DataFrame(rel_errors_euler)

# Convergence table
convergence_comparison_df = DataFrame(
    dt = [r[:dt] for r in results_exponential],
    exp_error_pct = [r[:obj_rel_error_pct] for r in rel_errors_exp],
    euler_error_pct = [r[:obj_rel_error_pct] for r in rel_errors_euler]
)

show(transform(convergence_comparison_df, names(convergence_comparison_df,
end

```

6×3 DataFrame			
Row	dt	exp_error_pct	euler_error_pct
	Float64	Float64	Float64
1	2.0	22.0589	1.5812
2	1.0	11.6362	0.7131
3	0.5	5.7315	0.3145
4	0.25	2.6017	0.1343
5	0.1	0.6581	0.0335
6	0.05	0.0	0.0

Plots

